



Numba: Expanding Capabilities to Power a New Generation of Python Libraries

Siu Kwan Lam, Sr Software Engineer, Anaconda



What is Numba?

Numba is a **just-in-time** **type-specializing** compiler
function

Demo



“What is Numba” demo key points

- **@jit** and **@njit** decorator
- **nopython=True** enforces native code compilation
- **type specialization** optimizes based on types
- **SIMD-vector** instructions



Slow to fast loops

- Does not support all python syntax
- Optimizes loops and array ops best
- “FORTRAN-style python”

Demo



“Slow to fast loop” demo key points

- Pure-python loop is inefficient
- LLVM enables loop-unroll, loop-vectorization
- Intel SVML
 - LLVM patch by Intel
 - `conda install -c numba icc_rt`
- **parallel=True** for auto-parallelization
 - Contributed by Intel

Expanding capabilities



Key features:

- @jit, @njit
- @vectorize, @guvectorize
 - cpu/gpu portable ufunc
- @cuda.jit, @rocm.jit
 - low level GPU kernel



Auto-parallelization era


In 2017 IntelLab contributed

ParallelAccelerator

@jit(parallel=True)

@stencil

prange



Library extension era

Since 2018


New extension libraries

- defines new containers

General purpose features

- list, dictionary, unicode
- try-except

Extensible compiler



OSS Projects that extend Numba



<https://github.com/scikit-hep/awkward-1.0>

“one of the most widely pip-installed packages for particle physics”

```
import awkward1 as ak
```

```
import numba as nb
```

```
@nb.jit(nopython=True)
```

```
def run(array):
```

```
    out = np.empty(len(array), np.float64)
```

```
    for i in range(len(array)):
```

```
        out[i] = array[i]["x"]
```

```
        for y in array[i]["y"]:
```

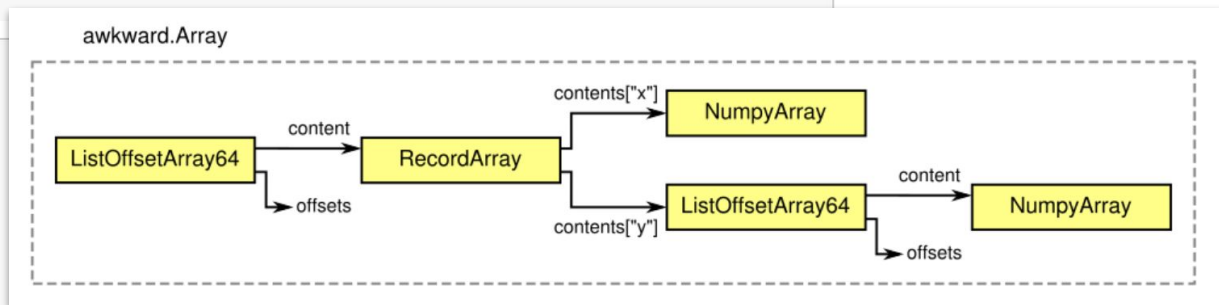
```
            out[i] += y
```

```
    return out
```

```
akarray = ak.Array([{"x": 100, "y": [1.1, 2.2]}, {"x": 200, "y": []}, {"x": 300, "y": [3.3]}])
```

```
# Works for the layout nodes, but not the high-level ak.Array wrapper yet.
```

```
run(akarray.layout)
```



reference: <https://github.com/scikit-hep/awkward-1.0/blob/master/docs-jupyter/2020-01-22-numba-demo-EVALUATED.ipynb>

Intel Scalable DataFrame Container (SDC)

```
import pandas as pd
from numba import njit, prange

# Dataset for analysis
FNAME = "employees.csv"

# This function gets compiled by Numba* and multi-threaded
@njit(parallel=True)
def get_analyzed_data():
    df = pd.read_csv(FNAME)
    s_bonus = pd.Series(df['Bonus %'])
    s_first_name = pd.Series(df['First Name'])

    # Use explicit loop to compute the mean. It will be compiled as parallel loop
    m = 0.0
    for i in prange(s_bonus.size):
        m += s_bonus.values[i]
    m /= s_bonus.size

    names = s_first_name.sort_values()
    return m, names

# Printing names and their average bonus percent
mean_bonus, sorted_first_names = get_analyzed_data()
print(sorted_first_names)
print('Average Bonus %:', mean_bonus)
```

```
$ python ./basic_workflow_parallel.py
7      ALEXANDER
4      CHRISTOPHER
0      EMILY
2      ISAAC
8      JOSEPH
9      JOSEPH
5      MIA
1      NOAH
6      OLIVIA
3      NaN
dtype: object
Average Bonus %: 11.204399999999998
```

Reference "Basic workflow in parallel" example:
<https://intelpython.github.io/sdc-doc/latest/examples.html>

Intel SDC Pipeline

- Extends Numba's pipeline
- Define custom passes

```
146 class SDCPipeline(numba.core.compiler.CompilerBase):
147     """SDC compiler pipeline
148     """
149
150     def define_pipelines(self):
151         name = 'sdc_extention_pipeline_distributed'
152         pm = DefaultPassBuilder.define_nopython_pipeline(self.state)
153
154         add_pass_before(pm, InlinePass, InlineClosureLikes)
155         pm.add_pass_after(HiFramesPass, InlinePass)
156         pm.add_pass_after(DataFramePass, AnnotateTypes)
157         pm.add_pass_after(PostprocessorPass, AnnotateTypes)
158         pm.add_pass_after(HiFramesTypedPass, DataFramePass)
159         pm.add_pass_after(DistributedPass, ParforPass)
160         pm.finalize()
161
162         return [pm]
163
164
165 @register_pass(mutates_CFG=True, analysis_only=False)
166 class ParforSeqPass(FunctionPass):
167     _name = "sdc_extention_parfor_seq_pass"
168
169     def __init__(self):
170         pass
171
172     def run_pass(self, state):
173         numba.parfors.parfor.lower_parfor_sequential(
174             state.typingctx, state.func_ir, state.typemap, state.calltypes)
175
176         return True
```

Extending with `@numba.extending.overload`

```
150 @overload(hq.heappush)
151 def heappush(heap, item):
152     assert_heap_type(heap)                                Typing logic
153     assert_item_type_consistent_with_heap_type(heap, item)
154
155     def hq_heappush_impl(heap, item):                      Codegen logic
156         heap.append(item)
157         _siftdown(heap, 0, len(heap) - 1)
158
159     return hq_heappush_impl
```


- Using setuptools entry points

```
138  setup(name = "awkward1",  
...  
211      entry_points = {  
212          "numba_extensions": ["init = awkward1._connect._numba:register"]  
213      },
```

from <https://github.com/scikit-hep/awkward-1.0/blob/master/setup.py>

Conclusion



Why libraries are using Numba?

- Native code performance
 - Without rewriting in pre-compiled languages
- Auto-parallelization
 - Leverage multicores CPU with ease
- Runtime codegen → easy to distribute
 - No need to ship binaries for every platform
- From numerical computing to data-analytics



Supported Platforms

- X86: Windows, Linux, Mac (64-bit only)
- ARMv7, AARCH64, PPC64: Linux
- Python ≥ 3.6
- GPU: CUDA, ROCM
- Distribute wheels and conda packages



Resources and examples

Demo notebooks (per release since v0.41):

<https://github.com/numba/numba-examples/tree/master/notebooks>

Doc & mailing list: <http://numba.pydata.org/>

Github: <https://github.com/numba/numba>

Gitter: <https://gitter.im/numba/numba>

Feel free to ask general questions on mailing list or Gitter, and open Github issues on specific problems.



Thank You!

Siu Kwan Lam
Github: @sklam

