

# Introduction to Spark using Python

[orlando.karam@gmail.com](mailto:orlando.karam@gmail.com)

# Administrivia

- If you have a question, ask right away
- We're being recorded

# What is Spark

- Distributed data processing framework
  - Distributed – runs in several machines
    - We get more RAM, more processing power
  - Data processing
    - Read and process data
- Based on Resilient Distributed Datasets
- Used for 'big data' processing
  - 'big' is something that doesn't fit on normal machines
    - Changes as machines become more powerful

# Resilient Distributed Datasets

- Imagine a big set of objects, and how we can distribute/parallelize
- We can divide in slices and keep each slice in a different node; this is the basic idea of an RDD
  - Values are computed only when needed
  - To guarantee fault-tolerance, we also keep info about how we calculated each slice, so we can re-generate it if a node fails
  - We can hint to keep in cache, or even save on disk
- Immutable ! not designed for read/write
  - instead, transform an existing one into a new one
- It is basically a huge list
  - But distributed over many computers

# Shared Spark Variables

- Broadcast variables
  - copy is kept at each node
- Accumulators
  - you can only add; main node can read

# Functional programming in python

- A lot of these concepts are already in python
  - But python community tends to promote loops
- Functional tools in python
  - map
  - filter
  - reduce
  - lambda
  - Itertools
    - Chain, flatmap

# Map in Python

- Python supports the map operation, over any list
- We apply an operation to each element of a list, return a new list with the results
  - `a=[1,2,3]`
  - `def add1(x): return x+1`
  - `map(add1,a) => [2,3,4]`
- We usually do this with a for loop, this is a slightly different way of thinking

# Filter

- Select only certain elements from a list
- Example:
  - `a=[1,2,3,4]`
  - `def isOdd(x): return x%2==1;`
  - `filter(isOdd,a) => [1,3]`



# reduce in python

- Applies a function to all pairs of elements of a list; returns ONE value, not a list
- Example:
  - `a=[1,2,3,4]`
  - `def add(x,y): return x+y`
  - `reduce(add,a) => 10`
    - `add(1,add(2,add(3,4)))`
- Better for functions that are commutative and associative, so order doesn't matter

# lambdas

- When doing map/reduce/filter, we end up with many tiny functions
- Lambdas allow us to define a function as a value, without giving it a name
- example: `lambda x: x+1`
  - Can only have one expression
  - do not write return
  - I put parenthesis around it, usually not needed by syntax
- `(lambda x: x+1)(3) => 4`
- `map(lambda x: x+1, [1,2,3])=> [2,3,4]`

# Exercises

- `(lambda x: 2*x)(3) => ?`
- `map(lambda x: 2*x, [1,2,3]) =>`
- `map(lambda t: t[0], [ (1,2), (3,4), (5,6) ] ) =>`
- `reduce(lambda x,y: x+y, [1,2,3]) =>`
- `reduce(lambda x,y: x+y, map(lambda t: t[0], [ (1,2), (3,4), (5,6) ] ))=>`

# More exercises

- Given
  - `a=[ (1,2), (3,4), (5,6)]`
- Write an expression to get only the second elements of each tuple
- Write an expression to get the sum of the second elements
- Write an expression to get the sum of the odd first elements

# Flatmap

- Sometimes we end up with a list of lists, and we want a 'flat' list
- Many functional programming languages (and Spark) provide a function called flatMap, which flattens such a list
- Example:
  - `Map(lambda t:range(t[0],t[1]), [ (1,5), (7,10)])` # returns list of lists
- `Itertools.chain` maps a list of iterables into a flat list
  - And so enables us to define our own flatmap

# Now let's do those with Spark

- Start the spark shell
  - run pyspark

# Creating RDDs in Spark

- All spark commands operate on RDDs (think big distributed list)
- You can use `sc.parallelize` to go from list to RDD
- Later we will see how to read from files
- Many commands are lazy (they don't actually compute the results until you need them)
- In pySpark, `sc` represents your `SparkContext`

# Simple example

- `list1=sc.parallelize( range(1,1000) )`
- `list2=list1.map(lambda x: x*10) # notice lazy`
- `list2.reduce(lambda x,y: x+y)`
- `list2.filter(lambda x: x%100==0).collect()`



# Transformations vs Actions

- We divide RDD methods into two kinds:
  - Transformations
    - return another RDD
    - are not really performed until an action is called (lazy)
  - Actions
    - return a value other than an RDD
    - are performed immediately

# Some RDD methods

- Transformations

- `.map( f )` – returns a new RDD applying `f` to each element
- `.filter( f )` – returns a new RDD containing elements that satisfy `f`
- `.flatMap(f)` – returns a ‘flattened’ list

- Actions

- `.reduce( f )` – returns a value reducing RDD elements with `f`
- `.take( n )` – returns `n` items from the RDD
- `.collect()` – returns all elements as a list
- `.sum()` - sum of (numeric) elements of an RDD
  - `max,min,mean ...`

# More examples

- `rdd1=sc.parallelize( range(1,100) )`
- `rdd1.map(lambda x: x*x).sum()`
- `rdd1.filter(lambda x: x%2==0).take(5)`

# Exercises

1. Get an RDD with number 1 to 10
2. Get all the elements in that RDD which are divisible by 3
3. Get the product of the elements in 2

# Reading files

- `sc.textFile(urlOrPath,minPartitions,useUnicode=True)`
  - Returns an rdd of strings (one per line)
  - Can read from many files, using wildcards (\*)
  - Can read from hdfs, ...
  - We normally use map right after and split/parse the lines
- Example:
  - `people=sc.textFile("../data/people.txt")`
  - `people=sc.textFile("../data/people.txt").map(lambda x: x.split('\t'))`

# Tuples and ReduceByKey

- Many times we want to group elements first, and then calculate values for each group
- In spark, we operate on tuples, <Key,Value> and we normally use reduceByKey to perform a reduce on the elements of each group

# People example/Exercises

- We have a people.txt file with following schema:
  - Name | Gender | Age | Favorite Language
- We can load with:
  - `people=sc.textFile("../data/people.txt").map(lambda x: x.split('\t'))`
- Find number of people by gender
  - first get tuples like: ('M',1),('F',1) ... then reduce by key
  - `people.map(lambda t: (t[1],1)).reduceByKey(lambda x,y:x+y).collect()`
- Let's find number of people by favorite programming language
- Example: youngest person per gender
  - `people.map(lambda t: (t[3],int(t[2]))).reduceByKey(lambda x,y:min(x,y)).collect()`

# More people exercises

- Get number of people with age 40+
  - Using filter
  - Using map and reduceByKey to produce two groups <40, 40+



# Person example with objects

- Using tuples for everything is ... *ok*, but sometimes we want nicer schema
  - We can use regular python objects
  - We still need to use tuples for joins, reduceByKey, since they operate on tuples
  - Can use x.name x.age etc which makes it slightly easier

# Person class

```
class Person:
    def parse(self,line):
        fields=line.split('\t')
        self.name=fields[0]
        self.gender=fields[1]
        self.age=int(fields[2])
        self.favorite_language=fields[3]
        return self

    def __repr__(self):
        return "Person( %s, gender=%s, %d years old, likes %s)"%
            (self.name,self.gender,self.age,self.favorite_language)
```

- `people=sc.textFile("../../data/people.txt").map(Person().parse)`

# Sending programs within shell

- You can use extra parameters to include python (or java) programs in your shell
  - --py-files (and list of files, separated with spaces)
    - Can use .py, .zip, .egg
  - --jars to include java jars
  - --packages, -- repositories to include maven packages (java)
  - --files to include arbitrary files in home folder of executor
- Get out of pyspark
  - Ctrl-D
- Run it again, including person.py in your --py-files

# Person with Objects

- Number of people by gender
  - `people.map(lambda t: (t.gender,1)).reduceByKey(lambda x,y:x+y).collect()`
- Let's do number of people by programming language
- Youngest person by gender
  - `people.map(lambda t: (t.gender,t.age)).reduceByKey(lambda x,y:min(x,y))`

# More people exercises

- Get number of people with age 40+
  - Using filter
  - Using map and reduceByKey to produce two groups <40, 40+
- Get age of oldest person, by programming language

# Sales example

- Sales: Day | StoreId | ProductId | QtySold
- Load:
  - `sales=sc.textFile("sales-data/sales_*.txt").map(lambda x: x.split('\t'))`
- now sales is an rdd of arrays corresponding to the fields
  - but each field is a string
- Total quantity of products sold:
  - `sales.map(lambda x: int(x[3])).sum()`

# Grouping RDDs again

- Work on RDDs of *pairs*,  $\langle key, value \rangle$
- `.reduceByKey(func)`
  - groups based on the key
  - reduce values in each group using the passed function
    - function is same way as reduce
  - produces RDD  $\langle key, result \rangle$

# Example

- `sales_by_store=sales.map( lambda t : (t[1], int(t[3])))`
- `sales_by_store.reduceByKey(lambda t1,t2: t1+t2).collect()`



# Exercises

- Calculate the sales for each day
- Calculate the total sales for each day for store 1
- Calculate the total sales for each product

# Joins

- Joins allow us to combine 2 different RDDs
  - Each RDD is of the form  $\langle K, V \rangle$  (key and value)
  - Result is of the form  $\langle K, \langle V1, V2 \rangle \rangle$  (notice the nesting)
  - Joins only on equal keys (equijoin from db)
  - Also have leftOuterJoin, rightOuterJoin and fullOuterJoin
  - And cartesian, if you want the cartesian product, and other kinds of joins, but this is potentially very slow

# Simple join example

```
states=[
  ("AL", "Alabama"),
  ("AK", "Alaska"),
  ("AR", "Arizona")
]; # apologies to the other 47 ...

populations=[
  ("AL", 4779736),
  ("AK", 710231),
  ("AR", 6392017)
]; # according to 2010 census, from Wikipedia

states_rdd=sc.parallelize(states)
populations_rdd=sc.parallelize(populations)

states_rdd.join(populations_rdd);
```

# Sales and Objects

- Two other files, one for Products one for Stores
  - Classes: Store, Product, SaleRow, with parse method
- `base_path="../data/sales"`
- `Sales_schema.py`

```
stores=sc.textFile(base_path+"stores*.txt").map(lambda x:sales_schema.Store().parse(x))
products=sc.textFile(base_path+"products.txt").map(lambda x:sales_schema.Product().parse(x))

sales=sc.textFile(base_path+"sales_*.txt").map(lambda x:sales_schema.SaleRow().parse(x))
```

# Sales examples (with objects)

- `sales_by_day=sales.map(lambda x : (x.day,x.quantity)).reduceByKey(lambda x,y:x+y)`
- `sales_by_store=sales.map(lambda x : (x.store_id,x.quantity) ).reduceByKey(lambda x,y:x+y)`
- Now let's do sales by product
- Get products with category stuff

# Sales and Joins (with objects)

- sales\_by\_store\_joined=
  - sales\_by\_store.join(stores.map(lambda x: (x.id,x.name)))
- Now let's do it with products

# Other joins

- Outer joins
  - Include the keys
  - .leftOuterJoin, .rightOuterJoin, .fullOuterJoin
- Cartesian Product
  - .cartesian

# Writing spark applications

- Need to obtain a SparkContext
  - from pyspark import SparkContext, SparkConf
  - conf = SparkConf().setAppName(appName) # appName not needed, but ...
  - sc = SparkContext(conf=conf)
- Everything is the same after that !!
- You probably want to save your data ... 😊
  - saveAsTextFile
  - Remember part files 😞
  - Can save in other ways



# Other functions

- Sample(withReplacement, fraction, seed)
- Union, intersection, distinct
- Coalesce, repartition
- aggregateByKey
- groupByKey
- repartitionAndSortWithinPartitions
- mapPartitions, mapPartitionsWithIndex

# New DataTable functionality

- A datatable is like an RDD but with schema information
  - Like a table in SQL, or datatable in pandas
  - Generic objects, know their fields
  - Datatable knows all its columns
  - All 'rows' are of the same kind (but there are nulls, and arrays etc)
- We need to either read from places with schemas, or add schema info
- We specify queries on them (similar to RDD, or through SQL), but there's a query optimizer
  - Slightly harder to do general aggregates
- Much smaller python tax !

# Person datatable example

- Easiest way to get data with schema is from a 'json' file
  - Each line is a json object
  - { "field": "value", ... }
- Need to use sqlCtx
  - `people=sqlCtx.jsonFile("../../data/people1.json")`
- Notice how each element is a Row, knows its fields
- `.show()` displays in nice way (first 20 by default)

# Datatable

- `.select` – like map, can use strings or columns
  - `people.select("name",people.age+1).show()`
- `.filter` – filter certain rows
  - `people.filter(people.age>30)`
- `.show` – display nicely
- Pandas syntax for filter
  - `people[people.gender=='F']`
- `GroupBy` returns a grouped RDD
  - `people.groupBy(people.gender).count()`
- `Join`

# Group By

- .GroupBy creates a grouped RDD
  - Can specify several fields
  - Still need to specify aggregates
- Aggregates
  - Count, sum, ...
- Can specify several with agg

# SQL

- Need to register the tables with the context
  - `people.registerTempTable("people")`
- Then can use `.sql` to do sql queries
  - `sqlCtx.sql("select name, age FROM people").show()`
  - `sqlCtx.sql("select gender,avg(age) AS Av FROM people GROUP BY gender")`

# Performance considerations

- Spark in python is slower than in scala due to translation
  - Spark processes are running in JVM
  - Need to send objects back and forth between jvm and python
- Databricks avoids this translation, it all lives in JVM
  - Until last step to client 😊
- Databricks can optimize better
  - But you lose some control
- Shuffling (join/reduce) is more expensive
  - Partitioning can help some

# RDD Performance

- RDD is:
  - Lineage
    - Set of Partitions/splits
    - List of dependencies on parent RDDs
    - Function to compute each partition given its parents
  - Optimized execution
    - Partitioner – which objects go on which partitions
      - Partitioning can help when shuffling
    - Preferred location for each partition



# Execution

- Your program
- Spark driver (master)
  - Keeps track of RDD graph
  - Scheduler
  - Block tracker
  - Shuffle tracker
- Spark executors
  - Task threads
  - Block manager